

Edge Computing

Lecture 08: Edge & Cloud: Middleware

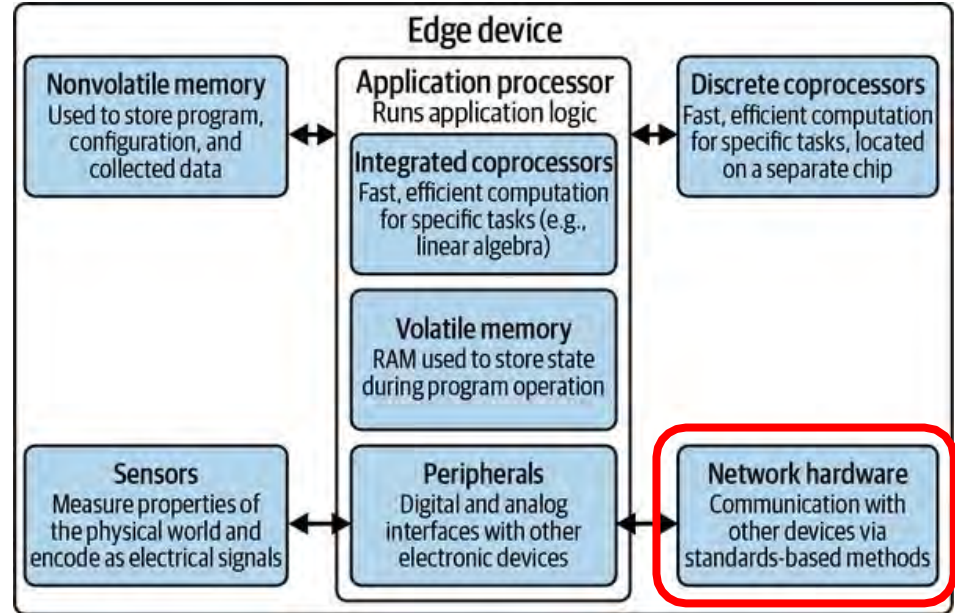
Paper sources

- Conference programs
 - MLSys: Machine Learning Systems, e.g. [MLSys'23](#)
 - Mobisys: Mobile Systems and Applications, e.g. [Mobisys'23](#)
 - CVPR: Computer Vision and Pattern Recognition, e.g. [CVPR'23](#)
 - SEC: Symposium on Edge Computing, e.g. [SEC'23](#)
- [Google scholar](#)
 - Search by topic, e.g. drone perception, smart agriculture
 - Search by name, e.g. [Bill Dally](#), [MLSys board member and steering committee](#)
- Source file: UCR Library
- Is it a good paper?
 - Top tier conference? [CSRankings](#)
 - High citation?
 - Github source code with many stars?

Recap

Hardware

- CPU
- Memory
- Cache
- RISC vs CISC
- Special accelerators
- Sensors



Agenda

Networking with others

- Middleware: What and why
- Design goals
- Examples

System Architect

- A person who designs hardware, software, or networking applications and services of a specified type for a business or other organization
- Edge computing: a designer's perspective


Designing an autonomous ride-sharing service

- What data? How does the data flow?



Designing an autonomous ride-sharing service

- What data? How does the data flow?



G
A
S

TCM

Transmission control module

PCM

Powertrain control module


ECM

Engine control module



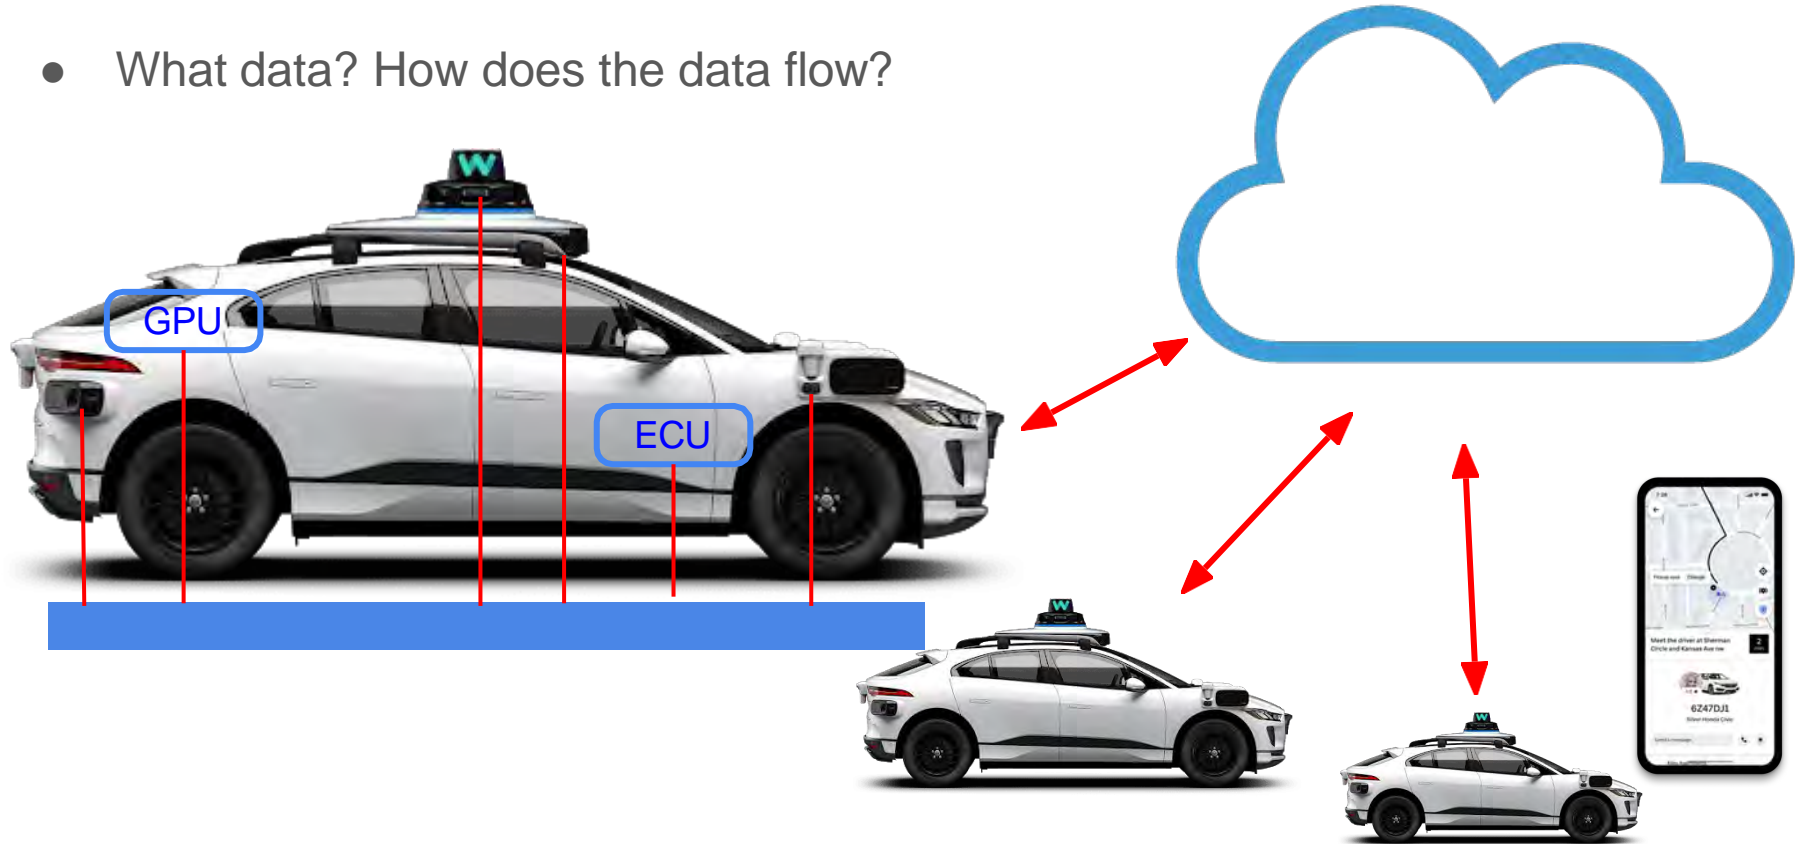
Designing an autonomous ride-sharing service

- What data? How does the data flow?



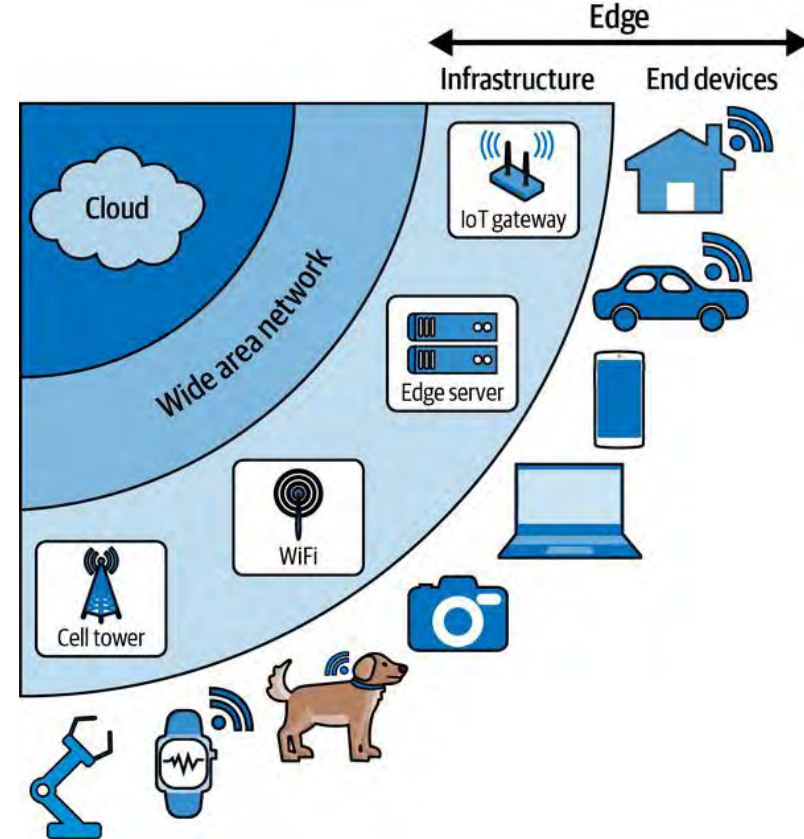
Designing an autonomous ride-sharing service

- What data? How does the data flow?




Middleware

- Broadly defined:
 - Software between OS and App.
- **Hidden translation layer**, enabling **communication and data management** for distributed applications.
- In edge computing context, a bridge
 - Among edge devices
 - Between edge and cloud



Middleware Design Goals

- Ad-hoc Discovery
 - Automated communication channel setup
 - Dynamically join or leave the channel
 - Data flow integrity and security



Middleware Design Goals

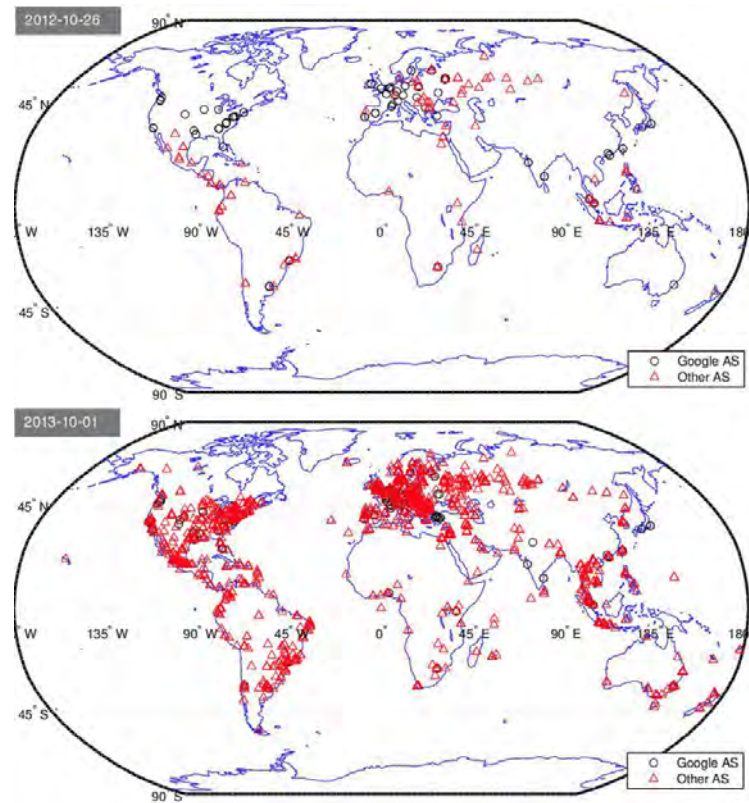
- Run-time Execution
 - Remotely execute programs (on edge devices or cloud) from a central hub
 - Deploy (code download)
 - Execution, start / stop service
 - Update (code)
 - Delivery of results



Google colab

Middleware Design Goals

- Minimal Task Disruption
 - Anticipate and handle disruptions
 - Reliability of middleware, via replica, backup, etc.
 - Resilience to mobility, network disconnections, memory overflow, etc.

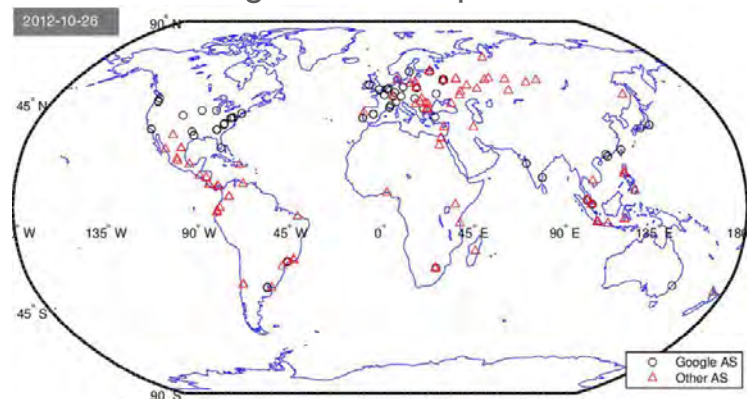


Google's CDN expansion

Middleware Design Goals

- Operation Overhead
 - Bandwidth
 - Latency
 - Compute
 - Energy

Google's CDN expansion



Main underwater cables




Middleware Design Goals

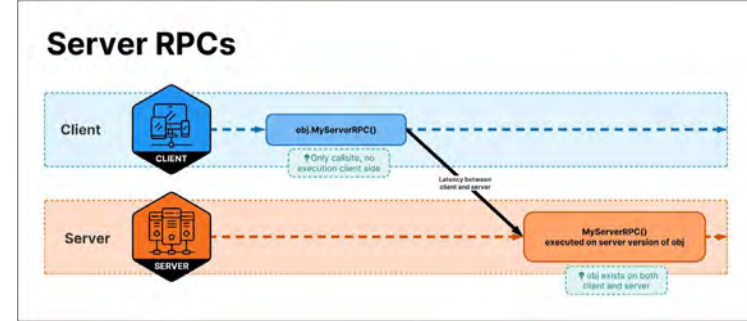
- Context-aware adaptive design
 - Hardware
 - Power budget
 - Network condition
 - User activity
 - Dynamics



Design Goal & Solutions

- Ad-hoc Discovery
 - Run-time Execution
 - Minimal Task Disruption
 - Operation Overhead
 - Context-aware adaptive design
- 
- Probing & selection of participants
 - RPCs, Virtualization
 - Replica, distributed system
 - Resource management, optimization
 - Context monitoring and prediction


Remote Procedure Calls (RPC)



- An object and its methods (or a procedure and its parameters) are “remoted” such that an invocation of a procedure or method can happen across a network separation.
- A local proxy (“stub”) mimic the interface of the remote object and its methods
- An application (client) make a call for remote execution (server) via the stub as if the call was executed locally.
- Tightly-coupled interfaces:
 - each application know the details of how other applications wants to communicate
- Synchronous vs Asynchronous?

Message-Oriented Middleware (MOM)

- Passing of data between applications using a communication channel that carries self-contained units of information (messages).
- Messages are usually sent and received asynchronously.
- Loosely coupled:
 - MOM API can finish send and receive, clients can be heterogeneous (python vs c++).
 - Clients don't need to know about each other's protocol/existence
- Asynchronous design (mostly)



Pub / Sub

- Publish / Subscribe
- Messaging pattern where publishers use message distribution nodes (brokers) where subscribers can retrieve relevant message.

Publish/ Subscribe Pattern



Pub / Sub

- Pros

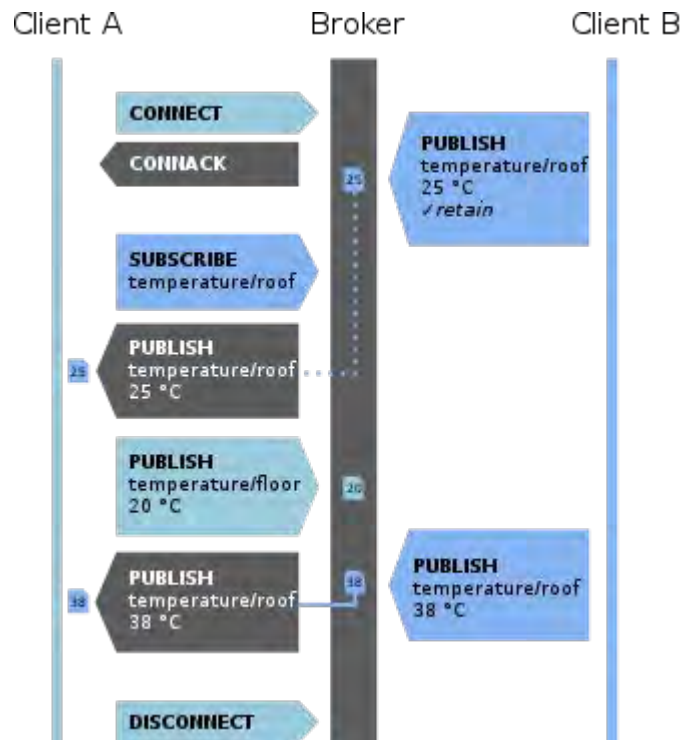

- Flexible, dynamic join (subscribe) and remove (unsubscribe)
- Loosely coupled, publishers and subscribers don't need to know about each other
- Scalable design

- Cons

- Reliability: single node of failure (broker/brokers)
- Network saturation
- Latency

MQTT

- Message Queue Telemetry Transport
 - A pub / sub network protocol
 - Four types of messages
 - Connect
 - Disconnect
 - Publish
 - Subscribe



MQTT Quality of Service

- At most once - (fire and forget)
 - the message is sent only once and the client and broker take no additional steps to acknowledge delivery.
- At least once - (acknowledged delivery)
 - the message is re-tried by the sender multiple times until acknowledgement is received.
- Exactly once - (assured delivery)
 - the sender and receiver engage in a two-level handshake to ensure only one copy of the message is received.

Example: Paho-MQTT + Mosquitto (broker)

- Install mosquitto broker: [Download | Eclipse Mosquitto](#)
- Install paho-mqtt: [paho-mqtt · PyPI](#)
- Start mosquitto
- Run the test script on the right
- Publish a message under topic “test”
 - `client.publish("test", "Hello world!")`

```
import paho.mqtt.client as mqtt

def on_connect(client, userdata, flags, rc):
    print("Connected with result code "+str(rc))
    client.subscribe("test")

def on_message(client, userdata, msg):
    print(msg.topic+" "+str(msg.payload))

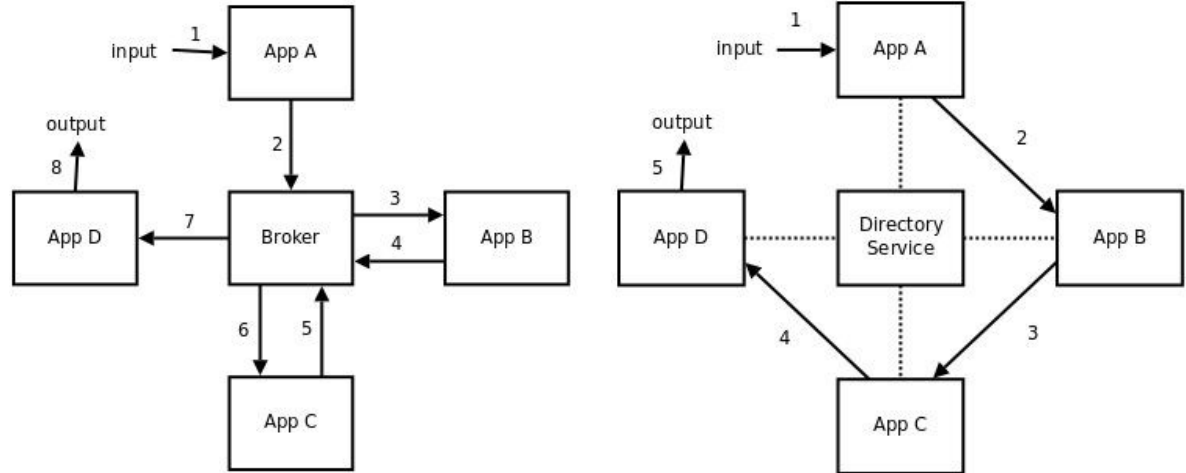
client = mqtt.Client()
client.on_connect = on_connect
client.on_message = on_message

client.connect("localhost", 1883, 60)

client.loop_start()
from IPython import embed; embed()
```

Brokerless MOM: Example: ZeroMQ or ZMQ

- Lightweight
 - Decouple broker's two function: pub / sub register, data transfer
 - Use directory service for pub / sub registration
 - Let clients handle data transfer
- N-to-N patterns
 - Fan-out
 - Pub-sub
 - Task-distribution
 - Request-reply



Other Examples:

1. [AWS IoT for the Edge](#)
2. [Azure Stack Edge](#)
3. [Kaa IoT Platform](#)
4. [Apache Kafka](#)
5. [ThingsBoard](#)
6. [openHAB](#)
7. [Eclipse Kura](#)
8. [FOGLAMP](#)
9. [EdgeX Foundry](#)
10. [Apache Edgent](#)

Other Examples:

1. [AWS IoT for the Edge](#)
2. [Azure Stack Edge](#)
3. [Kaa IoT Platform](#)
4. [Apache Kafka](#)
5. [ThingsBoard](#)
6. [openHAB](#)
7. [Eclipse Kura](#)
8. [FOGLAMP](#)
9. [EdgeX Foundry](#)
10. [Apache Edgent](#)

[Quiz 000: Middleware Examples](#)

- What does it include? (programming API, network modeling, device management, all the main features)
- License type (open source, pay as you go, contract based, etc.)
- At what layer is it deployed (cloud, Edge Network, Fog, Holistic: all of them)
- Is it general purpose or for specific cases?
- Compatibility: Programming language, OS, device, cloud service, etc.

Summary

- Middleware
- Design goals
- Examples
 - MOMs, Pub/Sub, MQTT

Next Lectures

- Lab 3: client-server communication
- Lab 4: connecting to the cloud
 - Google cloud platform (GCP)
 - Amazon Web Service (AWS)
- Paper presentation next week!
- Ethnics & Security
- Edge Computing Research
 - Guest Lecture from CISL lab